

LIGHTWEIGHT FINITE ELEMENT MESH DATABASE IN JULIA*

PETR KRYSL†

Abstract. A simple, lightweight, and fast, package in the programming language Julia for managing finite element mesh data structures is presented. The key role in the design of the data structures is granted to the incidence relation. This concept has some interesting implications for the simplicity and efficiency of the implementation. The entire library has less than 500 executable lines. The low memory requirements are also notable. The user of the library is given power over the decisions which mesh entities should be represented explicitly in the data structures, and which of the topological relationships should be computed and stored. This enables a small memory footprint, yet affords a sufficiently rich topology description capability.

Key words. finite element, mesh, topology, data structure, incidence relation

AMS subject classifications. 68Q25, 68R10, 68U05

1. Introduction. A number of mesh data structures have been proposed in the literature [2, 5, 16, 20, 4]: radial-edge, winged, half-edge and half-face, entity-based, etc. Usually with the goal of accommodating richer representations of functions on meshes, and supporting complex topological queries. Alas, flexibility and power to support mesh adaptation tend to increase the complexity of the implementation, and speed is hard-won in such designs. A common feature of these approaches is the use of pointers to objects in memory [1]. One disadvantage is of course that even when the indexes are 32-bit, the pointers on current machines are typically 64-bit. Consequently the memory used for such data structures is not insignificant.

Hence, array-based structures geared towards efficient access, and parsimonious storage of static meshes, also find a receptive ground: STK [7] and MOAB [21, 6] are array-based mesh structures. An often-cited example is the mesh data structure implemented in FENiCS [14]. It seems also possible to include in this list the innovative and unusual Sieve [11], which in its high-performance incarnation is available as DMPlex [13].

The goal of this paper is to present a simple, lightweight, and fast, package for managing finite element mesh data structures [17] in the programming language Julia [22, 3]. There are one or two points which the readers may find of interest. The key role assigned to the incidence relation appears to be a novel idea, which has some interesting implications for the simplicity and efficiency of the implementation. The entire library has less than 500 executable lines. The low memory requirements are also of notice. The present library leaves to the user of the library the decisions on (a) of which of the of mesh entities of the four manifold dimensions (cells, faces, edges, and vertices) to represent explicitly in the data structures, and (b) which of the 12 topological relationships to compute and store. This is in contrast to the usual “take it or leave it” design. Also, we do not use pointers to objects in memory. In fact, we believe that a major factor contributing to the efficiency and simplicity of our library is that it is *not* object-oriented. The implementation is simple and easy to understand thanks to the Julia programming language [22, 3].

The paper is organized as follows: We present the essential ideas and concepts in Section 2, and we describe the basic objects and operations. Section 3 provides

*Submitted to the editors DATE.

Funding: This work was not directly supported by a grant or contract.

†University of California, San Diego, CA (pkrysl@ucsd.edu, <http://hogwarts.ucsd.edu/~pkrysl/>).

44 some experimental data points concerning the usability, flexibility, and costs of the
 45 representation. Discussion and conclusions round off the paper in Section 4.

46 **2. Description of meshes.** In finite element analysis there is no such thing as
 47 “the mesh”. Even the simplest finite element program will require two meshes: one
 48 for the evaluation of the integrals over the interior, and one for the evaluation of the
 49 boundary integrals. Complex finite element programs typically work with a *multitude*
 50 of meshes, depending on the requirements of the application. Super-convergent patch
 51 recovery, mixed methods, high-order finite element methods with degrees of freedom
 52 at the edges, faces, and interiors, in addition to the nodes [20], discontinuous and
 53 hybrid Galerkin methods [8], nodal integration methods [12], and so on, need access
 54 to mesh entities at various levels of mesh topology. The present mesh library provides
 55 enough support for these complex applications, as will be described below.

56 On the other hand, many basic forms of the finite element method will require
 57 only the connectivity, enumerating for each element its nodes (i.e. a single downward
 58 adjacency). If that is so, for efficiency reasons there’s no point in constructing and
 59 storing additional topological information when it isn’t used. Hence, the present
 60 library can also attend to the needs of low complexity – low storage requirements
 61 cases.

62 In the next section we describe the basic objects¹ with which the library works:
 63 the shape descriptors, and the shape collections, the incidence relations, and the
 64 attributes. The reader may also find the Glossary in Appendix A to be of use.

65 **2.1. Shape descriptors, shapes, and shape collections.** We consider finite
 66 elements here to be *shapes*, such as line elements, triangles, hexahedra, etc. The
 67 shapes are classified according to their manifold dimension, so that we work with the
 68 usual vertices (0-dimensional manifolds), line segments (1-dimensional manifolds),
 69 triangles and quadrilaterals (2-dimensional manifolds), tetrahedra and hexahedra (3-
 70 dimensional manifolds).

71 The topology of an instance of the shape, which comprises information such as
 72 how many nodes are connected together, how many bounding facets there are and
 73 their definition, is described by *shape descriptors*. An example of a shape descriptor
 74 is provided in Figure 1 which shows the local topological description of a hexahedron
 75 shape. The encoding of the topological information into a shape descriptor allows for
 76 the functions constructing the incidence relations to work for any shape, no matter
 77 what the manifold dimension or order of the element.

78 The tables in Figure 1 introduce the concept of facets and ridges [15]: A *facet*
 79 is a bounding entity: faces for three-dimensional cells, edges for two-dimensional face
 80 elements, and vertices for one-dimensional line elements. A *ridge* is the “bounding
 81 entity of the bounding entity”. So edges are the ridges of the three-dimensional cells,
 82 and vertices are the ridges of the faces. Edges and vertices have no ridges. A good
 83 visual picture of facets and ridges may be provided by a finely cut diamond on the
 84 reader’s ring.

85 The shape itself is not oriented. However, the definition of the facets and ridges
 86 in terms of the vertices defines an inherent *orientation* (orientability) of the same.
 87 Therefore, our algorithms store the orientation of the uses of the facets and ridges in
 88 order to facilitate geometric queries. For instance, for the hexahedron the facets are

¹We wish to emphasize that we use the term object not in the sense of “object-oriented”. The programming language Julia [22, 3] itself is not object-oriented, and our implementation does not attempt graft itself upon the object-oriented tree.

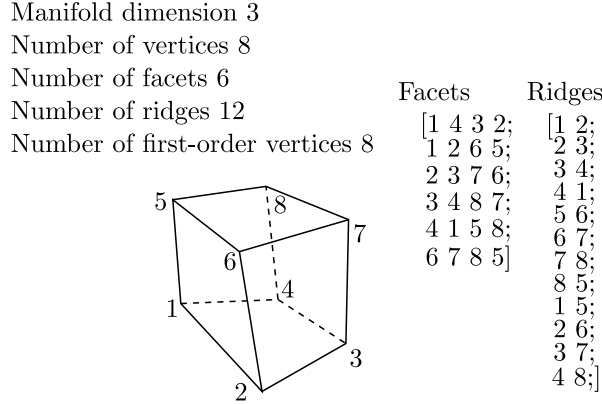


FIG. 1. The shape descriptor for an eight-node hexahedron element.

89 numbered so that when viewed from the outside of the hexahedron, each facet vertices
 90 are numbered counterclockwise. The ridges are numbered arbitrarily, as there is no
 91 intrinsic choice of numbering.

92 The shapes are considered in the form of collections: *Shape collections* are
 93 homogeneous collections of shapes. Collections of shapes do not hold any information
 94 about how the individual shapes are defined. That is the role of the incidence relations.
 95 The shape collections only provide information about the shape descriptor and the
 96 attributes of the shape collection, such as geometry (discussed below).

97 Finally, Figure 1 introduces the so-called *first-order vertices*. This concept
 98 is useful for applications of the library to high-order nodal elements, for instance.
 99 As introduced above, when computing relationships between three-dimensional cells
 100 and faces or between two dimensional cells (faces) and edges, it is useful to compute
 101 the orientation of the uses of the entity. As an example, a quadratic serendipity
 102 quadrilateral has eight vertices, but in order to figure out its orientation it is sufficient
 103 to refer to its four corner vertices. We call these the first-order vertices: they are the
 104 vertices of the first-order versions of the shapes.

105 **2.2. Incidence relation.** First, when do we consider entities of the mesh to be
 106 *incident*? An entity E of manifold dimension d_1 is considered to be incident on an
 107 entity e of manifold dimension $d_2 \leq d_1$, if e is contained in the *topological cover*
 108 of the entity E . So, as an example, a tetrahedron is incident on its faces, edges, and
 109 vertices. Due to our definition, a tetrahedron is also incident upon itself, but this last
 110 relation is hardly of any use, as is the case for all relations between entities $d_2 = d_1$.

111 Conversely, an entity e of manifold dimension $d_2 \leq d_1$ is incident on an entity E
 112 of manifold dimension d_1 if e belongs to E 's cover. So a vertex e is incident on all
 113 edges, faces, and cells that share it.

114 By *incidence relation* we mean here the relationship between two shape col-
 115 lections. We write

116 (2.1) (d_L, d_R)

117 where d_L is the manifold dimension of the shape collection on the left of the relation,
 118 and d_R is the manifold dimension of the shape collection on the right of the relation.
 119 The relationship can be understood as a function which takes as input a serial number
 120 of an entity from the shape collection on the *left* and produces as output a list of serial

TABLE 1

Table of incidence relations. Assuming that the initial mesh is three-dimensional, the first relationship to be established is the connectivity $(3, 0)$, as indicated by the box. The surface representation of the boundary, $(2, 0)$, would be a derived incidence relation. Other incidence relations may be subsequently computed as discussed in the text. MD = Manifold dimension.

MD	0	1	2	3
0	(0, 0)	(0, 1)	(0, 2)	(0, 3)
1	(1, 0)	(1, 1)	(1, 2)	(1, 3)
2	(2, 0)	(2, 1)	(2, 2)	(2, 3)
3	(3, 0)	(3, 1)	(3, 2)	(3, 3)

121 numbers of entities from the shape collection on the *right*, $i_L \rightarrow [j_{R,1}, \dots, j_{R,M}]$. Compare with Table 1 which lists the incidence relations that can be defined unambiguously between entities of the four manifold dimensions. The downward relationships are contained in the lower triangle of the matrix, moving from the bottom of the table upwards, and the upward relationships are listed top to bottom in the upper triangle of the matrix.

127 The relation $(0, 0)$ between two shape collections that consist of the same set of vertices, possibly in different order, is “trivial”: Vertex from the collection on the left is incident on itself in the collection on the right. This mapping may be a permutation, change of numbering. It is probably not worthwhile to actually create this relation, but it is included in Table 1 for completeness: it closes the computation of the skeleton (see below). The other relations between two shape collections (d, d) are included to complete the table, but the author is yet to find utility in these incidence relations.

135 Computational workflows typically start by creating a collection of d -dimensional shapes, where $d > 0$, such as a tetrahedral mesh produced by a mesh generator, and the collection of shapes is described by the **connectivity** (incidence relation) $(d, 0)$. This becomes the starting point for the computation of the required topological relations, as dictated by the needs of the particular finite element method (refer to Table 1). For definiteness, in the following we assume that we start with a three dimensional mesh (shown boxed in Table 1), so the basic data structure consists of the incidence relation $(3, 0)$. Should the initial mesh be two-dimensional, the table would be pruned by removing the fourth row and column.

144 **2.3. Derived Incidence Relations.** Here we address the issue of generating any of the other incidence relations of the table on the demand. For instance, the incidence relation $(2, 0)$ can be derived by application of the skeleton procedure to the incidence relation $(3, 0)$. Table 2 lists how the incidence relations in the rows and columns of the table are derived by listing the operation and its arguments. The relations on the diagonal for $d \geq 1$ are omitted, as they result by trivial permutation of the shape collection on the left into the shape collection of the right.

152 **2.4. skt: Skeleton.** The incidence relation $(2, 0)$ can be derived by application of the procedure “skeleton”. Repeated application of the skeleton will yield the relation $(1, 0)$, and finally $(0, 0)$. Note that at difference to other definitions of the incidence relation $(0, 0)$ (the paper of Logg comes to mind [14]) we consider this relation to be one-to-one, not one-to-many.

157 The skeleton procedure can be implemented in different ways. In our library we

TABLE 2

Operations to populate the table of incidence relations, starting from (3,0). *skt*=skeleton, *trp*=transpose, *bbf*= bounded-by facets, *bbr*= bounded-by ridges. *MD*= Manifold dimension.

MD	0	1	2	3
0	skt[(1,0)]	trp[(1,0)]	trp[(2,0)]	trp[(3,0)]
1	skt[(2,0)]	–	trp[(2,1)]	trp[(3,1)]
2	skt[(3,0)]	bbf[(2,0), (1,0), (0,1)]	–	trp[(3,2)]
3	(3,0)	bbr[(3,0), (1,0), (0,1)]	bbf[(3,0), (2,0), (0,2)]	–

158 use sorting of the connectivity of the entities of the skeleton as a two dimensional
 159 array in order to arrive at unique entities, eliminating duplicates (shared) entities.

160 **2.5. bbf: Bounded-by-facets.** The incidence relations (3,2) and (2,1) are ob-
 161 tained by the application of the “bounded-by-facets” procedure. In our implementa-
 162 tion the process draws upon three entity relations: the incidence of the mesh entities
 163 upon the vertices, and then bidirectional links between the facets and the vertices.

164 The facets are orientable. Therefore, our incidence relation stores signed entity
 165 numbers of the facets: when the facet use traverses the vertices of the facet in the
 166 same way in which the facet itself is stored, the orientation is positive (plus sign), and
 167 vice versa.

168 **2.6. bbr: Bounded-by-ridges.** The incidence relation (3,1) is obtained by the
 169 application of the “bounded-by-ridges” procedure. The process again draws upon
 170 three entity relations: the incidence of the cells upon the vertices, and then bidirec-
 171 tional links between the ridges and the vertices.

172 The ridges are orientable. Therefore, our incidence relation stores signed entity
 173 numbers of the ridges: when the ridge use traverses the vertices of the ridge in the
 174 same way in which the ridge itself is stored, the orientation is positive (plus sign),
 175 and vice versa.

176 As an aside, it would also be possible to generate the incidence relation (2,0) by
 177 the “bounded-by-ridges” procedure. It is of course also available by application of the
 178 skeleton procedure from the relation (3,0).

179 **2.7. trp: Transpose.** All the incidence relations below the diagonal of the
 180 matrix of Table 2 yield lists of entities of fixed cardinality. For example, the number
 181 of faces, edges, and vertices for hexahedron is always 6, 12, 8 respectively. On the
 182 contrary, the relationships in the upper triangle of the matrix are always of variable
 183 cardinality. For example, the number of tetrahedra around an edge [i.e. the incidence
 184 relation (1,3)] depends very much upon which edge it is. All the relations above
 185 the diagonal are obtained from the relations below the diagonal by the “transpose”
 186 operation.

187 **2.8. Constructing the full “one-level” representation.** The full “one-level”
 188 representation (refer, for example, to [9]), namely the incidence relations downward
 189 (3,2), (2,1), (1,0), and upward (0,1), (1,2), and (2,3) can be constructed by our
 190 library from the input (3,0) using the sequence of operations

$$\begin{aligned}
& (2, 0) = \text{skt}[(3, 0)] \\
& (0, 2) = \text{trp}[(2, 0)] \\
& (3, 2) = \text{bbf}[(3, 0), (2, 0), (0, 2)] \\
191 \quad (2.2) \quad & (1, 0) = \text{skt}[(2, 0)] \\
& (0, 1) = \text{trp}[(1, 0)] \\
& (2, 1) = \text{bbf}[(2, 0), (1, 0), (0, 1)] \\
& (1, 2) = \text{trp}[(2, 1)] \\
& (2, 3) = \text{trp}[(3, 2)]
\end{aligned}$$

192 **2.9. Incidence relations on the diagonal:** (d, d) . Logg [14] defines the inci-
193 dence relations (d, d) , where $d > 1$, as being one-to-many. For instance, the relation
194 $(3, 3)$ in [14] consists of all three-dimensional cells that share a vertex with the cell
195 on the left. Such incidence relations do not fit our definition of incidence of Section 2.2.
196 Even if we extend the definition of what we mean by “incident”, there are problems.
197 The definition of such a relation is not unique: In addition to the collections of shapes
198 on the left and on the right, it needs to refer to a connecting shape to make sense,
199 and hence it doesn’t fit Table 2. For instance, the relationship between faces, $(2, 2)$,
200 needs to state through which shape the incidence occurs: is it through a common
201 vertex? Is it through a common edge? Similarly, for cells the incidence relationship
202 $(3, 3)$ will be different for the incidences that follow from a common vertex, from a
203 common edge, or from a common face. This is one of the reasons we keep in Table 2
204 only the incidence relation $(0, 0)$. It fits the definition of incidence, and it is needed
205 as a closure of the skeleton operation.

206 **2.10. Mesh. Meshes** are understood here simply as incidence relations. At
207 the starting point of a computation, initial meshes are defined by the *connectivity*
208 of the finite elements and the finite element nodes, i.e. the incidence relation $(d, 0)$,
209 where $d \geq 0$, linking a d -dimensional shape to a collection of vertices as shapes in the
210 form of 0-dimensional manifolds. Any other mesh can be derived by the operations
211 of Table 2.

212 **2.10.1. Attributes.** At a minimum, the geometry of the mesh needs to be de-
213 fined by specifying the locations of the vertices. In our library we handle this data as
214 attributes of the shape collections. So the locations of the vertices are an attribute of
215 the shape collection of the vertices.

216 **2.11. Implementation notes.** Most mesh databases in current use favor the
217 storage of entity identifiers as 32-bit integers. This allows for substantial ranges of
218 approximately 2 billion positive and 2 billion negative identifiers (which may be useful
219 when storing orientation together with the serial number). If this is not enough, the
220 identifiers may be stored as 64-bit integers. This practically doubles the requisite
221 memory, but considerably expands the range. The present library accommodates
222 storage of the incidence relations with both and either integer types not only in the
223 same library, but also in the same running program: such is the magic of generic
224 programming as implemented in Julia [3] that in the same running program some
225 of the incidence relations may be stored as 32-bit integers while others are stored as
226 64-bit integers. This mixing is entirely transparent to the user. To get this to work
227 does not require anything beyond specifying parametric types.

228 As outlined above, the incidence relations below the diagonal differ from the inci-
229 dence relations above the diagonal by being of fixed cardinality. The implementation

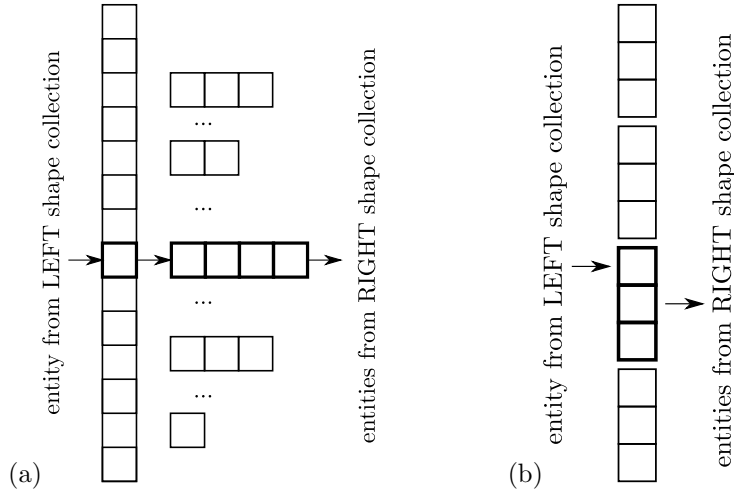


FIG. 2. (a) Storage of variable-cardinality vector of vectors on the left. (b) Storage of fixed-cardinality vector of vectors on the right.

230 in Julia can take proper advantage of this fact while maintaining a single interface
 231 to the incidence relations. The incidence relation is stored as a vector of vectors.
 232 Variable-cardinality incidence relations (above the diagonal of the matrix of Table 2)
 233 are stored as shown on the left. This is not as efficient as storing a fixed-cardinality
 234 vector of vectors: If all the vectors stored in the master vector are of fixed size, the
 235 package `StaticArrays` [10] can be used to enable operations on vectors that can be
 236 stored on the stack and that can be in-lined in a vector of vectors as shown in Figure 2.
 237 In the fixed-length case, each incidence vector is stored contiguously within one big
 238 array (on the right of the figure). Clearly, this saves memory as no storage of pointers
 239 for an indirection is needed.

240 The efficient storage of the fixed-cardinality vector of factors is enabled by Julia
 241 compiler’s ability to reason about the code, producing optimized implementation that
 242 can take advantage of any information that is known at compile time. At the same
 243 time, the programmer sees a uniform interface to the vector of vectors. This is the
 244 complete definition of the type of the incidence relation in our library:

```

245
246 struct IncRel{LEFT<:AbsShapeDesc, RIGHT<:AbsShapeDesc, T}
247     left::ShapeColl{LEFT} # left shape coll. (L, .)
248     right::ShapeColl{RIGHT} # right shape coll. (., R)
249     _v::Vector{T} # vec. of vec.s: shape num.s
250     name::String # name of the inc. relation
251 end
252

```

253 For instance, to find out how many entities in the shape collection on the right are
 254 linked to the j -th entity in the shape collection on the left we use the definition of
 255 the function

```

256
257 nentities(ir::IncRel, j) = length(ir._v[j])
258

```

259 Clearly, this function does not distinguish between fixed-cardinality and variable-
 260 cardinality vector of vectors.

261 **3. Results.** The computations described below were implemented in the Julia
 262 programming language [22, 3]. The mesh-topology library is implemented as the
 263 `MeshCore.jl` Julia package [17], and the computations referred to in this paper are
 264 available to the reader as part of the package `PaperMeshTopo.jl` [18].

265 An interesting comparison of the memory usage for the data structures can be
 266 gleaned from Figure 6 of [9]. The mesh is unfortunately not available directly, it is
 267 only known that it consists of 100,000 tetrahedra. Hence in the present system we
 268 simply generate a tetrahedral mesh of approximately 102,000 elements and compare
 269 the resulting storage requirements.

270 In Figure 3 we compare with the following systems: The MDS database of [9]
 271 was the full array representation. MDS-RED referred to as the “reduced array” was
 272 the element-to-vertex representation, both using 32-bit indices. Both MDS versions
 273 were storing vertex coordinates, geometric model classification, and coordinates of
 274 vertices. The MOAB database included element-to-vertex downward and upward
 275 adjacency. Apparently only element connectivities and vertices were stored in STK.
 276 All of these data bases stored 32-bit indices.

277 First our database was constructed to hold all of the incidence relations that
 278 correspond to the *full one-level storage* of MDS. That is we computed and stored
 279 the $(3, 2)$, $(2, 1)$, $(1, 0)$ and $(0, 1)$, $(1, 2)$, and $(2, 3)$ incidence relations. “MeshCore
 280 32” refers to this structure with indices stored as 32-bit integers, and “MeshCore
 281 64” refers to the equivalent topology structure with indices stored as 64-bit integers.
 282 When we store this information in 32-bit integers, we use only 68% of the memory
 283 compared to the MDS.

284 Next, our database was constructed to hold the incidence relations that corre-
 285 spond to the MOAB database with element-to-vertex downward and upward adja-
 286 cency. “MeshCore D/U” refers to this structure with indices stored as 32-bit integers.
 287 Hence we use only 39% of the storage of MOAB, and 61% of the storage for the
 288 MDS-RED.

289 Finally, a third data structure using our library, “MeshCore D”, stores only the
 290 $(3, 0)$ incidence relation. The storage requirement is an order of magnitude smaller
 291 than MDS, and amounts to around five times less memory than MOAB. It may not
 292 be an appropriate comparison in situations requiring more voluminous topological
 293 information, but if the finite element program has no use for the additional incidence
 294 relations, there’s no point in storing them, and a mesh storage scheme that can avoid
 295 this cost can win big. Our design can freely choose which incidence relations to store,
 296 and therefore we have fine control over the amount of stored information. That is
 297 the advantage of the structure of the data not being committed to by the design: the
 298 amount of information to be stored is left up to the user.

299 **4. Conclusions.** The library `MeshCore.jl` implements a storage model for meshes
 300 composed of common shapes such as triangles and quadrilaterals, tetrahedra and hex-
 301 ahedra. All incidence relations (sometimes known as adjacencies) that are commonly
 302 encountered in the literature can be produced by the library, which implements the
 303 four operations (skeleton, bounded-by-facets, bounded-by-ridges, and transpose) that
 304 can derive for instance the full one-level downward adjacencies (or downward *and*
 305 upward adjacencies, if desired). We avoid hardwiring the definition of the topological
 306 model in the implementation, at difference to common mesh databases. Our sep-
 307 aration of the data model and the implementation allows for a nimble and flexible
 308 computation of just the incidence relations that are actually needed. Consequently,
 309 the library is very conservative in terms of memory consumption.




FIG. 3. Comparison of the required memory to store various data structures. Legends are discussed in the text.

310 Importantly, we also avoid the use of pointers to memory, which is typical with
 311 object-oriented mesh databases [1]. Hence we avoid the penalty associated with stor-
 312 ing pointers at 64-bits, which is the norm on current computer architectures. Our
 313 Julia implementation stores the database in contiguous arrays whenever possible, and
 314 transparently switches to vector of vectors for variable-length data.

315 The current limitations include:

- 316 • The data structures may allow for adaptivity, but the current implementation
 317 of the library is static. At least in the sense that if the mesh changed, the
 318 incidence relations could be recalculated, but not in an incremental fashion.
- 319 • Homogeneous meshes are implemented. Mixed-shape meshes appear feasible,
 320 but have not been implemented yet.
- 321 • Support for non-manifold geometries is possible, but so far no effort was
 322 expended to reach this goal.
- 323 • No consideration has been given at this point to an extension for distributed
 324 databases for parallel computations.

325 The implementation in the Julia language produces code that can be at the same
 326 time flexible, powerful, and concise — the entire library has only around 490 exe-
 327 cutable lines, and with copious comments it clocks in at around 1000 lines. This may
 328 be contrasted with for instance the current version of MOAB which consists of two
 329 orders of magnitude larger number of lines of code. Of course, MOAB is much more
 330 powerful (it provides import/export, mesh modification, parallel execution). But the
 331 point could be made that that leaves open some room at the other end of the spec-
 332 trum: something flexible, easy to understand, and small in footprint. We believe that
 333 our library fits in that opening quite well.

334 An interesting opportunity for considerably expanding the usefulness of this li-
 335 brary has been identified by Rypl [19]: due to the generic form of the library, it

336 is suitable for operating on four-dimensional (for instance time space) meshes. The
 337 needed modification entails the addition of a shape descriptor for the four-dimensional
 338 cell. The three-dimensional cell would then become a facet, and the faces would be-
 339 come ridges. The tables of incidence relations would acquire a fifth row and column.

340 **Acknowledgments.** Continued support by the Office of Naval Research (pro-
 341 gram manager Michael J. Weise) is gratefully acknowledged. Daniel Ryp1 of the Czech
 342 Technical University in Prague is thanked for numerous thoughtful suggestions.

343 **Author contributions.** PK performed the work.

344 **Financial disclosure.** None reported.

345 **Conflict of interest.** The author declares no potential conflict of interests.

346

REFERENCES

- 347 [1] M. W. BEALL AND M. S. SHEPHARD, *A general topology-based mesh data structure*, In-
 348 ternational Journal for Numerical Methods in Engineering, 40 (1997), pp. 1573–1596,
 349 [https://doi.org/10.1002/\(sici\)1097-0207\(19970515\)40:9<1573::Aid-nme128>3.0.Co;2-9](https://doi.org/10.1002/(sici)1097-0207(19970515)40:9<1573::Aid-nme128>3.0.Co;2-9).
- 350 [2] L. L. BEGHINI, A. PEREIRA, R. ESPINHA, I. F. M. MENEZES, W. CELES, AND G. H. PAULINO,
 351 *An object-oriented framework for finite element analysis based on a compact topological*
 352 *data structure*, Advances in Engineering Software, 68 (2014), pp. 40–48, [https://doi.org/](https://doi.org/10.1016/j.advengsoft.2013.10.006)
 353 [10.1016/j.advengsoft.2013.10.006](https://doi.org/10.1016/j.advengsoft.2013.10.006).
- 354 [3] J. BEZANSON, A. EDELMAN, S. KARPINSKI, AND V. B. SHAH, *Julia: A fresh approach to numer-*
 355 *ical computing*, SIAM review, 59 (2017), pp. 65–98, <https://doi.org/10.1137/141000671>.
- 356 [4] W. CELES, G. H. PAULINO, AND R. ESPINHA, *A compact adjacency-based topological data struc-*
 357 *ture for finite element mesh representation*, International Journal for Numerical Methods
 358 in Engineering, 64 (2005), pp. 1529–1556, <https://doi.org/10.1002/nme.1440>.
- 359 [5] G. COMPERE, J. F. REMACLE, J. JANSSON, AND J. HOFFMAN, *A mesh adaptation framework*
 360 *for dealing with large deforming meshes*, International Journal for Numerical Methods in
 361 Engineering, 82 (2010), pp. 843–867, <https://doi.org/10.1002/nme.2788>.
- 362 [6] V. DYEDOV, N. RAY, D. EINSTEIN, X. M. JIAO, AND T. J. TAUTGES, *AHF: array-based half-*
 363 *facet data structure for mixed-dimensional and non-manifold meshes*, Engineering with
 364 Computers, 31 (2015), pp. 389–404, <https://doi.org/10.1007/s00366-014-0378-6>.
- 365 [7] H. C. EDWARDS, A. B. WILLIAMS, G. D. SJAARDEMA, D. G. BAUR, AND W. K. COCHRAN,
 366 *Sierra toolkit computational mesh conceptual model*, Sandia National Laboratories SAND
 367 series, technical report, SAND2010-1192, University of Minnesota, Albuquerque, NM, 2010.
- 368 [8] M. S. FABIEN, M. G. KNEPLEY, R. T. MILLS, AND B. M. RIVIERE, *Manycore parallel comput-*
 369 *ing for a hybridizable discontinuous Galerkin nested multigrid method*, Siam Journal on
 370 Scientific Computing, 41 (2019), pp. C73–C96, <https://doi.org/10.1137/17m1128903>.
- 371 [9] D. IBANEZ AND M. S. SHEPHARD, *Modifiable array data structures for mesh topology*, Siam
 372 Journal on Scientific Computing, 39 (2017), pp. C144–C161, [https://doi.org/10.1137/](https://doi.org/10.1137/16m1063496)
 373 [16m1063496](https://doi.org/10.1137/16m1063496).
- 374 [10] JULIA ARRAYS DEVELOPMENT TEAM, *Statically sized arrays for Julia StaticArrays.jl*. <https://github.com/JuliaArrays/StaticArrays.jl>. Accessed 03/25/2020.
- 375 [11] M. G. KNEPLEY AND D. A. KARPEEV, *Mesh algorithms for PDE with Sieve I: Mesh distribution*,
 376 Scientific Programming, 17 (2009), pp. 215–230, <https://doi.org/10.1155/2009/948613>.
- 377 [12] P. KRYSL AND B. ZHU, *Locking-free continuum displacement finite elements with nodal inte-*
 378 *gration*, International Journal for Numerical Methods in Engineering, 76 (2008), pp. 1020–
 379 1043, <https://doi.org/10.1002/nme.2354>.
- 380 [13] M. LANGE, L. MITCHELL, M. G. KNEPLEY, AND G. J. GORMAN, *Efficient mesh management*
 381 *in Firedrake using PETSC DMPLEX*, Siam Journal on Scientific Computing, 38 (2016),
 382 pp. S143–S155, <https://doi.org/10.1137/15m1026092>.
- 383 [14] A. LOGG, *Efficient representation of computational meshes*, International Journal of Computa-
 384 tional Science and Engineering, 4 (2009), pp. 283–295, [https://doi.org/10.1504/ijcse.2009.](https://doi.org/10.1504/ijcse.2009.029164)
 385 [029164](https://doi.org/10.1504/ijcse.2009.029164).
- 386 [15] J. MATOUŠEK, *Lectures on Discrete Geometry*, Graduate Texts in Mathematics, Springer, New
 387 York, 1st ed., 2002.
- 388 [16] C. OLLIVIER-GOOCH, L. DIACHIN, M. S. SHEPHARD, T. TAUTGES, J. KRAFTCHECK, V. LEUNG,
 389 X. J. LUO, AND M. MILLER, *An interoperable, data-structure-neutral component for mesh*
 390

- 391 *query and manipulation*, Acm Transactions on Mathematical Software, 37 (2010), <https://doi.org/10.1145/1824801.1824807>.
- 392
- 393 [17] PETR KRYSL, *MeshCore: Lightweight Mesh Library in Julia*. [https://doi.org/10.5281/zenodo.](https://doi.org/10.5281/zenodo.3737253)
- 394 [3737253](https://doi.org/10.5281/zenodo.3737253), Accessed 03/23/2020.
- 395 [18] PETR KRYSL, *PaperMeshTopo: Computations for paper "Lightweight Mesh Library in Julia"*.
- 396 <https://github.com/PetrKryslUCSD/PaperMeshTopo.jl>, Accessed 03/23/2020.
- 397 [19] D. RYPL, *Oral communication*. 2020.
- 398 [20] E. S. SEOL AND M. S. SHEPHARD, *Efficient distributed mesh data structure for parallel au-*
- 399 *tomated adaptive analysis*, Engineering with Computers, 22 (2006), pp. 197–213, <https://doi.org/10.1007/s00366-006-0048-4>.
- 400 [21] T. J. TAUTGES, R. MEYERS, K. MERKLEY, C. STIMPSON, AND C. ERNST, *MOAB: A mesh-*
- 401 *oriented database*, Sandia National Laboratories SAND series, technical report, SAND2004-
- 402 1592, University of Minnesota, Albuquerque, NM, 2004.
- 403 [22] THE JULIA PROJECT, *The Julia programming language*. <https://julialang.org/>, Accessed
- 404 04/11/2019.
- 405

406 Appendix A. Glossary.

- 407 **Topological cover:** A cover of a set X is a collection of sets whose union contains
- 408 X as a subset.
- 409 **Shape:** Topological shape of any manifold dimension, 0 for vertices, 1 for edges, 2
- 410 for faces, and 3 for cells.
- 411 **Shape descriptor:** Description of the type of the shape, such as the number of
- 412 vertices, facets, ridges, and so on.
- 413 **Shape collection:** Set of shapes of a particular shape description.
- 414 **Facet:** Shape bounding another shape. A shape is bounded by facets: The facet is a
- 415 $d - 1$ -dimensional face of a d -dimensional entity.
- 416 **Facet use:** Facets are orientable. The incidence relation stores facet uses: when a
- 417 facet use orders the vertices in the same way (modulo circular shift) as the
- 418 referenced entity, the facet use is stored as a positive entity number; otherwise
- 419 it is stored as a negative entity number.
- 420 **Ridge:** Shape one manifold dimension lower than the facet. For instance a tetrahe-
- 421 dron is bounded by facets, which in turn are bounded by edges. These edges
- 422 are the "ridges" of the tetrahedron. The ridges can also be thought of as a
- 423 "leaky" bounding shapes of 3-D cells. The ridge is a $d - 2$ -dimensional face
- 424 of a d -dimensional entity.
- 425 **Ridge use:** Ridges are orientable. The incidence relation stores ridge uses: when a
- 426 ridge use orders the vertices in the same way (modulo circular shift) as the
- 427 referenced entity, the ridge use is stored as a positive entity number; otherwise
- 428 it is stored as a negative entity number.
- 429 **Incidence relation:** Map from one shape collection to another shape collection. For
- 430 instance, three-dimensional finite elements (cells) are typically linked to the
- 431 vertices by the incidence relation $(3, 0)$, i. e. for each tetrahedron the four
- 432 vertices are listed. Some incidence relations link a shape to a fixed number
- 433 of other shapes, other incidence relations are of variable arity. This is what
- 434 is usually understood as a "mesh".
- 435 **Incidence relation operations:** The operations defined in the library are: the
- 436 skeleton operation, the transpose operation, the bounded-by-facets opera-
- 437 tion, and the bounded-by-ridges operation. All topological relations between
- 438 the shapes of the four manifold dimensions that are uniquely defined can be
- 439 constructed using the sequence of these operations.
- 440 **Mesh topology:** The mesh topology can be understood as an incidence relation
- 441 between two shape collections.